

**A PROJECT REPORT**

**on**

**Toxic Comment Classification**

**In Partial Fulfilment of the Requirement for the Award of**

**BACHELOR'S DEGREE IN  
INFORMATION TECHNOLOGY**

**BY**

<b>Pragyesh Kumar Sinha</b>	1606034
<b>Shubham Gupta</b>	1606059
<b>Soumyajeet Addya</b>	1606066
<b>Srijan Mehrotra</b>	1606229
<b>Uditya Kumar</b>	1606233

**UNDER THE GUIDANCE OF**

**PROF. Ajay Kumar Jena**

**SCHOOL OF COMPUTER ENGINEERING  
KALINGA INSTITUTE OF INDUSTRIAL TECHNOLOGY  
BHUBANESWAR, ODISHA -751024**



# CERTIFICATE

This is certify that the project entitled

Toxic Comment Classification

submitted by

Pragyesh Kumar Sinha	1606034
Shubham Gupta	1606059
Soumyajeet Addya	1606066
Srijan Mehrotra	1606229
Uditya Kumar	1606233

is a record of bonafide work carried out by them, in the partial fulfilment of the requirement for the award of Degree of Bachelor of Engineering (Computer Science & Engineering OR Information Technology) at KIIT Deemed to be university, Bhubaneswar. This work is done during the year 2019-2020, under my guidance.

Date: 31 / 03 / 2020

(Prof. Ajay Kumar Jena)  
Project Guide

## **Acknowledgements**

We are profoundly grateful to Prof. Ajay Kumar Jena for his expert guidance and continuous encouragement throughout to see that this project rights its target since its commencement to its completion.

Pragyesh Kumar Sinha  
Shubham Gupta  
Soumyajeet Addya  
Srijan Mehrotra  
Uditya Kumar

## ABSTRACT


The Internet is used globally all around the world and it offers a lot of anonymity to anyone who wants to express their views on any particular subject. The anonymity that the internet offers can be harmful because it allows people to say whatever they want without any one being accountable for it. Personal attacks, bullying are really common in the comment area of YouTube, Twitter, etc.

As a result Toxic Comment Analysis has become a really vast area for people, industry and scientists to perform research on. In this project, we used Wikipedia's talk page edits to train a different multi-level classifier that detects different levels of toxicity. We will first clean the data to remove links, IP-addresses, image-links, CSS code, wiki templates, and also process the internet's shorthand. We will then use Tf-idf Vectorizer and Fasttext to vectorize the cleaned corpus to produce the vector of the text. Which we will pass through different deep learning models: Artificial Neural Networks (ANN), Convolutional Neural Networks (CNN), and Recurrent Neural Networks (RNN) with LSTM to predict the six different types of toxicity: toxic, severe toxic, obscene, threat (eg threat to kill someone) insult, identity hate (hate against a specific group of people). Finally we will use Bi-directional Gated Recurrent Unit with a Convolutional layer, to get the final accuracy and F1-score.

**Keywords:** NLP, Neural Networks, CNN, RNN, LSTM, Tfifdf, Bi-GRU.

# Contents

1.	Introduction	2
2.	Literature Survey	3
3.	Software Requirement Specification	4-5
3.1.	Tools	4
3.1.1.	Jupyter Notebook	4
3.2.	Framework	4-5
3.2.1.	Numpy	4
3.2.2.	Pandas	4
3.2.3.	Matplotlib	4
3.2.4.	Keras	5
3.2.5.	Tfidf Vectorizer	5
4.	Requirement Analysis	6
4.1.	Vectorizer	6
4.2.	Neural Networks	6
5.	Project Planning	7
5.1.	Data Analysis	7
5.2.	Data Cleaning	7
5.3.	Data Pre-processing	7
5.4.	Model Training and Testing	7
6.	Implementation	8-14
6.1.	Analysis	8-9
6.2.	Cleaning	9-11
6.3.	Pre-Processing	11
6.4.	Training and Testing	12-14
7.	Screenshots	15-20
8.	Conclusion and Future Report	21
8.1.	Conclusion	21
8.2.	Future Report	21
9.	Reference	22



# List of Figures

8.1: Multiple Toxicity Bar graph	10
8.2: Toxicity Heat-Map	10
8.3: Clean Text Function	11
8.4: Architecture of NN	12
8.5: Architecture of CNN	12
8.6: Architecture of RNN	12
9.1: Basic import of NN	13
9.2: Basic import of CNN	13
9.3: Basic import of RNN	14
9.4: Bar Plot for Clean and Toxic	14
9.5: CNN Accuracy and Loss	15 – 16
9.6: RNN Accuracy and Loss	16 – 17
9.7: NN Accuracy and Losses	17

# Chapter 1

## Introduction

**Toxic Comments** are those comments that have the capability to spark some kind of controversy between - two people, two groups of people, or a person and a group of people. It generally consists of rude things about a specific person or a group of person in general. It can also be a radical idea that can't be supported by a lawful society or that doesn't bring about any positive changes in society. The Internet has become a common grounds for preaching such toxic comments, simply because of the anonymity that the internet grants.

People on the Internet can say anything they want and get away with it without having any kind of accountability for what they say. And this kind of toxic comments is not only harmful to the proper usage of the internet but also harmful for the people on the internet. People on the internet might get bullied, threatened or might face racism and bigotry for their identity, color, race, etc. This comment not only are a nuisance but also can be mentally taxing for people who face it. Thus for properly using the internet without facing any kind of toxic comments, those comments should be recognized and removed. It can also be a daunting task if proper techniques and tools are not used. In this project, we will use classification tools to recognize toxic comments.

**Classification** is the process or action of classifying something, in which things are put into categories. In this project, there are six types of toxic categories in which we can put a comment in toxic, severe\_toxic, threat, identity\_hate, obscene, insults. A comment can fall under one or many of these categories. If a comment doesn't fall under any of these categories then the comment must be free of any kind of toxicity and is clean. There are many different kinds of classification machine learning models that can be used to classify this: SVM, Naive Bayes, KNN, Decision Tree, Random forest, etc. In this project, we will use different kinds of Neural Networks for the classification of the comments.

**Neural Networks** are inspired by the networks of neurons present in our brain. A neural network is a connection of neurons or nodes. The connections between two neurons are modeled as weights which can be positive or negative, based on if we need excitatory or inhibitory connections. There are many different types of Neural Networks that are used for classification: Neural Networks, Convolutional Neural Networks, Recurrent Neural Networks with LSTM or GRU. We can also combine different models to make models such as GRU with Convolutional Networks, LSTM with Convolutional Networks, etc

# **Chapter 2**

## **Literature Survey**

Many previous research works have been done on Toxic Comment Classification. Many studies have used SLM like Linear Regression, SVM, NB to classify the toxicity of the comments, also many research papers try to see if there is any difference between classifying the data before preprocessing and after preprocessing. Many even try to clean the data, like remove links, ip addresses, etc, many don't.

In this project, learning from all those who researched it previously, we will be using NN to build the classifiers, we will also clean the models of not useful stuff like Links, IP addresses. And also pre-process it using three different pre-processor.

Besides classic binary classification tasks, related work examines different viewpoints of toxic languages, such as “racism” (Greevy and Smeaton, 2004; Waseem, 2016; Kwok and Wang, 2013) and “sexism” (Waseem and Hovy, 2016; Jha and Mamidi, 2017), or the hardness of toxicity (David-son et al., 2017; Sharma et al., 2018).

From this we got the inspiration to classify the data using Deep Learning Models instead of simple Binary Classification or Shallow Learning Models

# **Chapter 3**

## **Software Requirements Specification**

### **3.1 Tools**

#### **3.1.1 Jupyter notebook**

The Jupyter Notebook is an open source application that can run python code and display the output immediately. It also uses RELP which makes it flexible to use and really easy to make ML models and run ML projects.

### **3.2 Framework**

#### **3.2.1 NumPy**

NumPy is a library that supports arrays and many other numerical functions which are required during the pre-processing of the data. It can be used to manipulate large matrices which is really helpful for performing many different numerical functions

#### **3.2.2 Pandas**

Pandas is used to manipulate data in the data frame. It is really helpful in the data frame manipulation. It imports and exports the data in a data frame.

#### **3.2.3 Matplotlib**

Matplotlib is used to plot graphs. It is really helpful to analyse the data to get insight on the data with the help of graphs. We can plot a variety of graphs using this library, we can plot histogram, bar graph, box plot, violin plot, etc.

#### **3.2.4 Keras**

Keras is an open-source neural-network library written in Python. It uses the back-end framework of TensorFlow. Thus it is helpful in making Deep Learning models like NN, CNN, RNN, LSTM, GRU, etc. It makes the work of making the architecture of the model relatively easy.

#### **3.2.5 Tf-Idf Vectorizer**

TF-IDF is a process which assigns a number to the word in a text depending on the frequency of times the text appears on the comment and also how many times it appears on the documents. It vectorizes the data based on the frequency of times the word appears on the model.

# **Chapter 4**

## **Requirement Analysis**

### **4.1 Vectorization**

In Machine Learning the features are numerical data on which mathematical operations can be performed to predict the results. But in our project, the main data is text data. Thus the text data has to be converted to numerical data to predict toxicity. The task of converting the text data to numerical data is done by a vectorizer. Vectorizers are of many types: Tfifd Vectorizer, Countvectorizer, Google Word2Vec, FastText, etc.

### **4.2 Neural Networks**

Neural Networks work with the base concept that a layer of neurons gets numerical data from features which are added by multiplying and adding different weights. The result is then passed through a formula, eg: Sigmoid or Relu, etc. The result is passed to another layer of Neurons. Thus the process continues. In this project, we will use Neural Networks and variations of Neural Networks like:

1. CNN: Which firsts embeds the numerical data to make it more compact and then runs the similar concept of Basic Neural Networks.
2. RNN with LSTM: LSTM has a special property with the help of which it can remember the data that came before the current data, thus making the processing of the data more accurate. RNN is a sequence of neural network blocks that are linked to each other like a chain. Thus RNN is better at dealing with sequential data and since our data is textual data - and textual data is sequential - we can predict it will perform better than the Neural Networks or CNN.

# **Chapter 5**

## **Project Planning**

### **5.1 DATA ANALYSIS**

The data contains 150 K rows with 8 columns: ID, comment\_text, toxic, severer\_toxic, obscene, threat, insult, identity\_hate. The last six columns have values either 0 or 1 which denote if the comment text is toxic or not. The data is going to be explored to find correlations between unique words and toxicity and not on sentences and toxicity.

### **5.2 DATA CLEANING**

The comment text is going to be explored to find un-important text in the comments and this un-important text will be cleaned

### **5.3 DATA PREPROCESSING**

We can preprocess the data in three different ways: Vectorizing the data using Tfifdf Vectorizer. Using fast text word embedding of Wikipedia. Using badwords.csv predefined by Google.

### **5.4 MODEL TRAINING AND TESTING**


The four types of preprocessed data will be used to train and test three different deep learning models: Neural Networks, Convolutional Neural Networks and RNN under Long Short Term Memory, and Bi-GRU (using CNN).

# Chapter 6

## Implementation


### 6.1 ANALYSIS

The data was explored thoroughly to see if there are any correlations between no of words per comment and toxicity and it was plotted using matplotlib. The relationship between different types of toxicity was also found out and it was seen that every comment in severe\_toxic class was also toxic.



Multiple toxicity Bar Graph

We also found out the correlations between the different type of toxicity. Correlations between toxic-obscene, toxic-threat, and obscene-threat is high compared to other combinations. But since correlation is not causation, we can't conclude that if a comment is toxic it will also be obscene and vice-versa.



## 6.2 CLEANING

The comment text contains both lower and upper case letters, links, IP addresses, wiki tokens, CSS templates and short-hands. In this step we removed the IP addresses, wiki tokens, links and CSS templates, and then converted the whole text to lowercase and converted the short-hands to long-hands. This text cleaning is done using 'RE' library

### Code of cleaning:

```
def clean_text(cmnt_text, clean_wiki_tokens = True):
    cmnt_text = cmnt_text.lower()
    #removing links
    cmnt_text = re.sub(r"https?:\/\/(www\.)?[-a-zA-Z0-9@:%._\+~#=]{2,256}\.[a-z]{2,6}\b([-a-zA-Z0-9@:%_\+~#=?\&//=]*)", "", cmnt_text)
```

```
#removing IP addresses
cmnt_text =
re.sub(r"(25[0-5]|2[0-4][0-9]|01)?[0-9][0-9]?(.(25[0-5]|2[0-4][0-9]|01)?[0-9][0-9]?)){3}", "", cmnt_text)

if clean_wiki_tokens:
    #removing images
    cmnt_text = re.sub(r"image:[a-zA-Z0-9]*\.jpg", " ", cmnt_text)
    cmnt_text = re.sub(r"image:[a-zA-Z0-9]*\.png", " ", cmnt_text)
    cmnt_text = re.sub(r"image:[a-zA-Z0-9]*\.gif", " ", cmnt_text)
    cmnt_text = re.sub(r"image:[a-zA-Z0-9]*\.bmp", " ", cmnt_text)

    #removing CSS
    cmnt_text = re.sub(r"#([A-Fa-f0-9]{6}|[A-Fa-f0-9]{3})", " ", cmnt_text)
    cmnt_text = re.sub(r"\{[^{}]*\}", " ", cmnt_text)

    #removing templates
    cmnt_text = re.sub(r"\[?\\[user:\\*\\]", " ", cmnt_text)
    cmnt_text = re.sub(r"\[?\\[wikipedia:\\*\\]", " ", cmnt_text)
    cmnt_text = re.sub(r"\[?\\[special:\\*\\]", " ", cmnt_text)
    cmnt_text = re.sub(r"\[?\\[category:\\*\\]", " ", cmnt_text)

    cmnt_text = re.sub(r"what's", "what is ", cmnt_text)
    cmnt_text = re.sub(r"\'s", " ", cmnt_text)
    cmnt_text = re.sub(r"\'ve", " have ", cmnt_text)
    cmnt_text = re.sub(r"can't", " cannot ", cmnt_text)
    cmnt_text = re.sub(r"\n't", " not ", cmnt_text)
    cmnt_text = re.sub(r"\i'm", " i am ", cmnt_text)
    cmnt_text = re.sub(r"\m", " i am ", cmnt_text)
    cmnt_text = re.sub(r"\re", " are ", cmnt_text)
    cmnt_text = re.sub(r"\d", " would ", cmnt_text)
    cmnt_text = re.sub(r"\ll", " will ", cmnt_text)
    cmnt_text = re.sub(r"\,", " ", cmnt_text)
    cmnt_text = re.sub(r"\.", " ", cmnt_text)
    cmnt_text = re.sub(r"\!", " ! ", cmnt_text)
    cmnt_text = re.sub(r"\?", " ? ", cmnt_text)
    cmnt_text = re.sub(r"\!", " ! ", cmnt_text)
    cmnt_text = re.sub(r"\\"", " ", cmnt_text)
    cmnt_text = re.sub(r"\^", " ^ ", cmnt_text)
    cmnt_text = re.sub(r"\+", " + ", cmnt_text)
    cmnt_text = re.sub(r"\-", " - ", cmnt_text)
    cmnt_text = re.sub(r"\=", " = ", cmnt_text)
    cmnt_text = re.sub(r"\\"", " ", cmnt_text)
    cmnt_text = re.sub(r"(\d+)(k)", r"\g<1>000", cmnt_text)
    cmnt_text = re.sub(r"\:", " : ", cmnt_text)
    cmnt_text = re.sub(r" e g ", " eg ", cmnt_text)
    cmnt_text = re.sub(r" b g ", " bg ", cmnt_text)
```

```

cmnt_text = re.sub(r" u s ", " american ", cmnt_text)
cmnt_text = re.sub(r"\0s", "0", cmnt_text)
cmnt_text = re.sub(r" 9 11 ", "911", cmnt_text)
cmnt_text = re.sub(r"e - mail", "email", cmnt_text)
cmnt_text = re.sub(r"j k", "jk", cmnt_text)
cmnt_text = re.sub(r"\s{2,}", " ", cmnt_text)
cmnt_text = re.sub(r"\n", " ", cmnt_text)

return(cmnt_text)

```

### 6.3 Preprocessing

The comment text needs to be converted in vectors which can be used to train the models. We primarily used three different types of pre-processing:

1. Vectorizing the data using TfIdf Vectorizer. It depends on the number of times a ‘term’ or word is present in a comment and how many times the ‘term’ occurs in the whole corpus of comments.
2. Fast-Text embedding, which is a predefined file which contains the vector for many different words. It is made by Facebook AI or FBAI.
3. Bad words. It is a pre-defined csv file with some of the most commonly used bad words. This csv file was used to Tokenize the words of the corpus.

### 6.4 Training and Testing

#### 1. Neutral Networks

The cleaned, Tf Idf Vectorizer data was used to train Neural Networks with 4 layers and 2 drop- outs.

#### Code of NN:

```

model = Sequential()
model.add(Dense(1000, input_dim = 500, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(250, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(32, activation='relu'))
model.add(Dense(num_class, activation='sigmoid'))
adam = optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-08, decay=0.0)
model.compile(loss='binary_crossentropy', optimizer=adam, metrics=['accuracy'])
model.summary()

```

## 2. Convolutional Neural Networks

The cleaned data pre- processed using Fast- Text was used to train CNN with 1 Embedding, 2 Convolution, 1 drop-out and 2 layers.

### Code of CNN:

```
print("training CNN ...")
model = Sequential()
model.add(Embedding(nb_words, embed_dim,
                   weights=[embedding_matrix], input_length=max_seq_len, trainable=False))
model.add(Conv1D(num_filters, 7, activation='relu', padding='same'))
model.add(MaxPooling1D(2))
model.add(Conv1D(num_filters, 7, activation='relu', padding='same'))
model.add(GlobalMaxPooling1D())
model.add(Dropout(0.5))
model.add(Dense(32, activation='relu', kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Dense(num_classes, activation='sigmoid')) #multi-label (k-hot encoding)

adam = optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-08, decay=0.0)
model.compile(loss='binary_crossentropy', optimizer=adam, metrics=['accuracy'])
model.summary()
```

## 3. Recurrent Neural Networks

The cleaned data pre- processed using BadWords.csv is used to train RNN with LSTM (Long Short Term Memory) with 1 Embedding and 4 dense layers.

### Code of RNN:

```
embedding_vecor_length = 300
model = Sequential()
model.add(Embedding(5000, embedding_vecor_length, input_length=X_train.shape[1]))
model.add(LSTM(256))
model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dense(6, activation='sigmoid'))

# Compiling Model using optimizer
opt = Adam(lr=1e-3)
model.compile(loss='binary_crossentropy', optimizer=opt, metrics=['accuracy'])
model.summary()
```

#### 4. Bi-GRU-CNN

The cleaned data is preprocessed using Fast Text and a bidirectional GRU using CNN layer is used to classify the data.

##### Code of Bi-GRU\_CNN:

```
from keras.optimizers import Adam, RMSprop
from keras.callbacks import EarlyStopping, ModelCheckpoint, LearningRateScheduler
from keras.layers import GRU, BatchNormalization, Conv1D, MaxPooling1D

file_path = "bi_gru_cnn.hdf5"
check_point = ModelCheckpoint(file_path, monitor = "val_loss", verbose = 1,
                               save_best_only = True, mode = "min")
ra_val = RocAucEvaluation(validation_data=(X_val, y_val), interval = 1)
early_stop = EarlyStopping(monitor = "val_loss", mode = "min", patience = 5)

def build_model(lr = 0.0, lr_d = 0.0, units = 0, dr = 0.0):
    inp = Input(shape = (max_len,))
    x = Embedding(max_features, embed_size, weights = [embedding_matrix], trainable =
False)(inp)
    x = SpatialDropout1D(dr)(x)

    x = Bidirectional(GRU(units, return_sequences = True))(x)
    x = Conv1D(64, kernel_size = 2, padding = "valid", kernel_initializer = "he_uniform")(x)
    avg_pool = GlobalAveragePooling1D()(x)
    max_pool = GlobalMaxPooling1D()(x)
    x = concatenate([avg_pool, max_pool])

    x = Dense(6, activation = "sigmoid")(x)
    model = Model(inputs = inp, outputs = x)
    model.compile(loss = "binary_crossentropy", optimizer = Adam(lr = lr, decay = lr_d),
metrics = ["accuracy"])
    model.summary()
    history = model.fit(X_train, y_train, batch_size = 128, epochs = 4, validation_data =
(X_val, y_val),
                           verbose = 1, callbacks = [check_point, early_stop])
    model = load_model(file_path)
    return model
```

# Chapter 7

## ScreenShots

### Basic Imports of NN

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, f1_score
import re, string

from keras import optimizers
from keras.models import Sequential
from keras import regularizers
from keras.layers import Dense, Activation, Dropout, Flatten
from keras.layers import Embedding, Conv1D, MaxPooling1D, GlobalMaxPooling1D
from keras.callbacks import EarlyStopping

from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer
from string import punctuation
from collections import defaultdict
import sys

Using TensorFlow backend.
```

### Basic Imports of CNN

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import keras
from keras import optimizers
from keras import backend as K
from keras import regularizers
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Flatten
from keras.layers import Embedding, Conv1D, MaxPooling1D, GlobalMaxPooling1D
from keras.utils import plot_model
from keras.preprocessing import sequence
from keras.preprocessing.text import Tokenizer
from keras.callbacks import EarlyStopping

from tqdm import tqdm
from nltk.corpus import stopwords
from nltk.tokenize import RegexpTokenizer
import os, re, csv, math, codecs

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, f1_score, recall_score

sns.set_style("whitegrid")
np.random.seed(0)

MAX_NB_WORDS = 100000
tokenizer = RegexpTokenizer(r'\w+')
stop_words = set(stopwords.words('english'))
stop_words.update(['.', ',', '"', "'", ':', ';', '(', ')', '[', ']', '{', '}'])

from subprocess import check_output
import warnings
warnings.filterwarnings('ignore')

Using TensorFlow backend.
```

## Basic Imports of RNN

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

import re

# Importing required libraries
from sklearn.model_selection import train_test_split
from nltk.corpus import stopwords

# keras
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.utils import to_categorical

# from keras.layers import Embedding

from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence

from keras.callbacks import EarlyStopping
from keras.optimizers import Adam

Using TensorFlow backend.
```

```
In [2]: def clean_text(cmnt_text, clean_wiki_tokens = True):
```


## Basic Imports of BI\_GRU\_CNN

```
In [1]: import time
start_time = time.time()
from sklearn.model_selection import train_test_split
import sys, os, re, csv, codecs, numpy as np, pandas as pd
np.random.seed(32)
os.environ["OMP_NUM_THREADS"] = "4"
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.layers import Dense, Input, LSTM, Embedding, Dropout, Activation, Conv1D
from keras.layers import Bidirectional, GlobalMaxPool1D, MaxPooling1D, Add, Flatten
from keras.layers import GlobalAveragePooling1D, GlobalMaxPooling1D, concatenate, SpatialDropout1D
from keras.models import Model, load_model
from keras import initializers, regularizers, constraints, optimizers, layers, callbacks
from keras import backend as K
from keras.engine import InputSpec, Layer



Using TensorFlow backend.
```

```
In [2]: def clean_text(cmnt_text, clean_wiki_tokens = True):
```


## Bar Plot for Clean and Toxic




## NN accuracy and loss




## CNN accuracy and loss



## RNN Accuracy and loss



RNN Loss



## Bi-GRU-CNN Accuracy and F1-Score

```
print("recall_score: ",recall_score(y_test[:, i], p[:, i]))
print('\n')

0
[[21165    510]
 [ 357   1904]]
f1_score:  0.8145454545454546
Accuracy:  0.9637784090909091
recall_score:  0.8421052631578947

1
[[23603     98]
 [ 133   110]]
f1_score:  0.4966139954853273
Accuracy:  0.9906834893048129
recall_score:  0.45267489711934156

2
[[22546    144]
 [ 256   990]]
f1_score:  0.8319327731092437
Accuracy:  0.9832887708534759
recall_score:  0.7945425361155698

3
[[23847    271]
 [ 38    24]]
f1_score:  0.4247787610619469
Accuracy:  0.9972844251336899
recall_score:  0.3870967741935484

4
[[22468    324]
 [ 225   919]]
f1_score:  0.770004189359028
Accuracy:  0.9770638368983957
recall_score:  0.8033216783216783

5
[[23723      8]
 [ 181    24]]
f1_score:  0.26253164556962028
Accuracy:  0.9921039438502673
recall_score:  0.11707317073170732
```

**Test Accuracy 98.39**

**Test f1\_score 59**

**Train Accuracy 98.42**

**Validation Accuracy 98.38**

In [ ]:

# **Chapter 8**

## **Conclusion**

In this project we saw that the deep learning models gave good results on a cleaned text data. Bi-GRU using a convolutional layer gave the best results with 98.4% accuracy. The second best result was given by RNN with 97% accuracy. CNN gave 96.6% accuracy. The Neural Network gave only 94.4% accuracy. The CNN and RNN gave 0 F1 score in threat, identity\_hate, whereas this wasn't the case in Bi-GRU and NN. From the Neural Network accuracy graph we can say that Neural Network overfit into the training data, that did not happen in Bi-GRU. Thus we can say Bi-GRU is the best model till now.

## **Future Work**

There is a lot of scope with future research/project work in this domain:

- We can use this model in a web extension which can be implemented on websites like YouTube and Twitter to detect and self-delete toxic comments.
- We can use Data Augmentation to make the data less imbalanced and get better accuracy, and can check what effect that has in the F1-score of CNN and RNN and if it increases the score of Bi-GRU.

# Reference

- [1] Xiang Zhang, Junbo Zhao, and Yann LeCun. 2015. Character- level convolutional networks for text classification. In Advances in neural information processing systems.
- [2] Charu C Aggarwal and ChengXiang Zhai. 2012. A survey of text classification algorithms. In Mining text data. Springer, 163–222.
- [3] Hossein Hosseini, Sreeram Kannan, Baosen Zhang, and Radha Poovendran. 2017. Deceiving Google’s Perspective API Built for Detecting Toxic Comments. arXiv preprint arXiv:1702.08138 (2017).
- [4] “Kaggle Toxic Classification Challenge.” [Online]. Available: [www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge](http://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge)
- [5] CrowdFlower, CrowdFlower Data Science Report, 8–9, 2016.
- [6] A. L. Maas, R. E. Daly, P. T. Pham, D. Huang, A. Y. Ng, and C. Potts, “Learning Word Vectors for Sentiment Analysis,” Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies, 142–150 (2011).
- [7] Fabio Del Vigna, Andrea Cimino, Felice Dell’Orletta, Marinella Petrocchi, and Maurizio Tesconi, Hate me, hate me not: Hate speech detection on facebook (2017).
- [8] Sanjana Sharma, Saksham Agrawal, and Manish Shrivastava, Degree based classification of harmful speech using twitter data (2018).
- [9] Steven Zimmerman, Udo Kruschwitz, and Chris Fox, Improving hate speech detection with deep learning ensembles (2018).
- [10] Mohammad, F., Is preprocessing of text really worth your time for online comment classification? (2018)
- [11] Conversational AI Research,  
<https://www.github.com/conversationai/unintended-ml-bias-analysis> (2018).